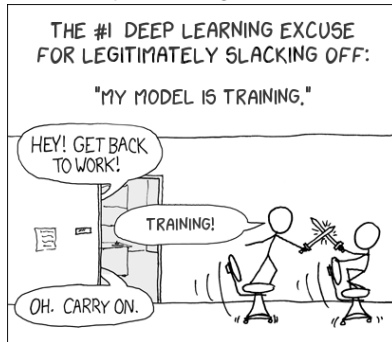


# EE3-25: Deep Learning

Krystian Mikołajczyk & Carlo Ciliberto

Department of Electrical and Electronic Engineering  
Imperial College London



## Course Information

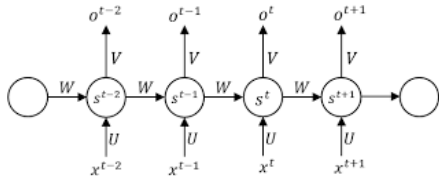
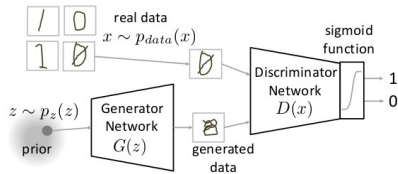
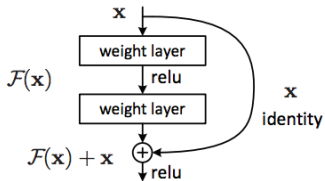
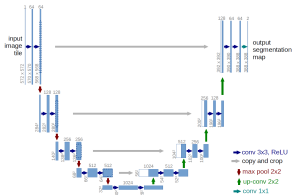
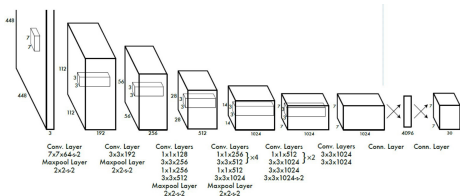
- Dr Krystian Mikolajczyk
  - Room 1015
  - Office hour: Friday 17:00pm-18:00pm
  - Email: [k.mikolajczyk@imperial.ac.uk](mailto:k.mikolajczyk@imperial.ac.uk)
- Dr Carlo Ciliberto
  - Room 1005
  - Office hour: Friday 17:00pm-18:00pm
  - Email: [c.ciliberto@imperial.ac.uk](mailto:c.ciliberto@imperial.ac.uk)
- GTAs
  - Axel Barroso Laguna
  - Adrian Lopez Rodriguez
  - Office hour: Tuesday 17:00pm-18:00pm

## Goal

- To introduce fundamental principles, theory and approaches for learning with deep neural networks.
- To offer practical course on implementing and experimenting with deep learning.
- Part of a course on machine learning and related topics:
  - EE3-10 (Autumn): Maths for Signals and Systems
  - EE3-23 (Autumn): Machine Learning
  - EE3-25 (Spring): Deep Learning
  - EE3-08 (Spring): Advance Signal Processing
  - EE4-68 (Autumn): Pattern Recognition
  - EE4-62 (Spring): Selected Topics in Computer Vision
  - EE4 (Spring): Final Year Project

## Goal

- Learn to apply different types of networks in various DL tasks



## Course Information

- Lectures: Friday 3-5pm (403)
  - Lecture notes available on the lecture day (at latest)
  - Book: *Deep Learning*, Ian Goodfellow, Yoshua Bengio & Aaron Courville, 2016. MIT Press, <http://www.deeplearningbook.org>
  - <https://towardsdatascience.com>
- Weekly practical exercises
  - Self studying on your PC
  - Colab - online python environment with Keras and TensorFlow backend
  - Lab sessions (Lab 305, Tuesday 17:00, 28/01, 4/02)
- Coursework (100%)
  - Work: online exercises from DL github, experiments and reports
  - Assessment: 2 page interim report 20%, Deadline: 13 Feb 2020 (23:59), via Blackboard
  - Assessment: 4 page final report 80%, Deadline: 19 March 2020 (23:59), via Blackboard

## Lectures by Dr Krystian Mikolajczyk

- Part 1: Introduction to deep learning
- Part 2: Convolutional Neural Networks (CNN)
- Part 3: Network Training
- Part 4: CNN architectures
- Part 5: Recurrent Neural Networks

- Part 6: Representation Learning and Autoencoders
- Part 7: GANs & friends
- Part 8: Metalearning
- Part 9: Reinforcement Learning I
- Part 10: Reinforcement Learning II

## Practical Experiments & Coursework

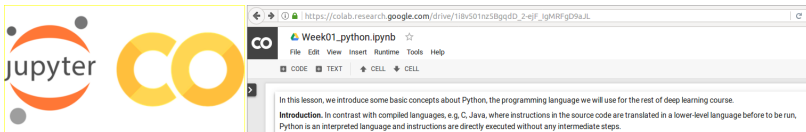
<https://github.com/MatchLab-Imperial/deep-learning-course>

- Week 2-3
  - Introduction to Python and some frameworks (NumPy, Pandas, etc..)
  - Introduction to Keras
- Week 4-6
  - Fundamentals of deep learning: handling different type of data (text, image, audio, etc), feedforward of artificial neural networks, introduction to last generation of CNN architectures (VGG, Inception, ResNet, UNets etc...)
- Week 7-10
  - Advanced deep learning: LSTM sequence modelling, Generative Adversarial Networks, Neural style transfer (CycleGan, Pix2Pix), Reinforcement Learning.



# Practical Experiments & Coursework

- Environment: Colaboratory <sup>\*1,2</sup>
  - Repository: <https://github.com/MatchLab-Imperial/deep-learning-course>
  - Jupyter notebook environment which requires not setup and supported from most major browsers, e.g, Chrome and Firefox.
  - Code is run in virtual machines with free GPU.
  - Files are stored securely in your own Google Drive account.
  - Supports developing Python applications using popular deep learning libraries, e.g, **Keras**, Tensorflow, Pytorch.



<sup>\*1</sup> <https://colab.research.google.com/notebooks/welcome.ipynb>

<sup>\*2</sup> <https://medium.com/deep-learning-turkey/google-colab-free-gpu-tutorial-e113627b9f5d>

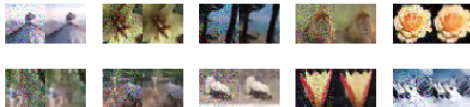
## Practical Experiments & Coursework

- Format: Jupyter <sup>\*3</sup>

- Notebooks are documents produced by the Jupyter Notebook Apps, e.g., Colaboratory, containing both python code and rich text elements (paragraph, equations, figures, links, etc. . . )

```
In [0]: N=5
start_val = 0# pick an element for the code to plot the following N*2 values
fig, axes = plt.subplots(N,N)
for row in range(N):
    for col in range(N):
        idx = start_val+row+N*col

        im = np.concatenate((np.clip(X_test_noise[idx], 0, 1), np.clip(pred[idx], 0, 1)), 1)
        axes[row,col].imshow(im)
        y_target = int(y_train[idx])
        axes[row,col].set_xticks([])
        axes[row,col].set_yticks([])
```



<sup>\*3</sup> [https://jupyter-notebook-beginner-guide.readthedocs.io/en/latest/what\\_is\\_jupyter.html](https://jupyter-notebook-beginner-guide.readthedocs.io/en/latest/what_is_jupyter.html)

## Practical Experiments & Coursework

- Deep Learning Framework: Keras <sup>\*4</sup>
  - modular, minimalist framework, especially good for beginners
  - along with Colab environment allows to set a neural network and start prototyping in no time.



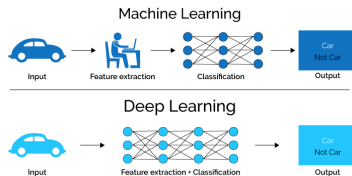
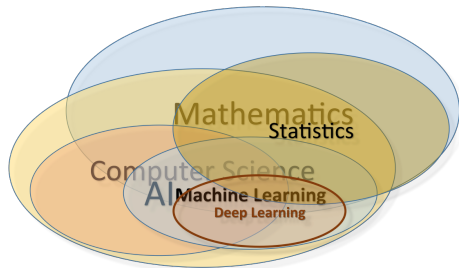
<sup>\*4</sup> <https://pypi.org/project/Keras/>

- BlackBoard Q&A forum
- Open Office hour, Lecturers and GTAs (humans)
- Emails - risk of getting missed, otherwise will be copied to Q&A forum anyway

# Deep Learning

# Deep Learning

- AI, Machine Intelligence
  - Intelligent agents with perception and actions to achieve goals
- Machine Learning
  - Ability to learn: Data  $\rightarrow$  Hypothesis
- Deep Learning
  - Ability to learn data representation (features) and predictors



Goal

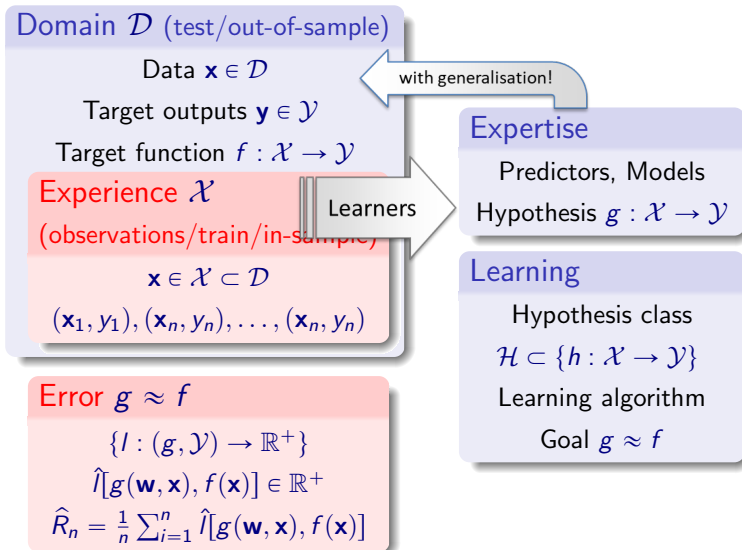
**Learning with generalisation**

### Machine Learning Revision

- Components of Learning
- ML Tasks
- Types of learning
- Types of data
- Learning setup
- Error/Loss Measures
- Perceptron
- Neural Networks
- Gradient descent
- Backpropagation
- Learning curves
- Regularization

## Components of learning

- Theory
  - ML problem formulation
  - Errors, loss and bounds
- Predictors and Learners
  - Linear and non linear
  - SVM
  - **Neural Networks**
- Learning frameworks
  - Supervised
  - Reinforcement learning





## Error Measures/Loss Functions

- How to quantify  $h \approx f$ ?
- Usually pointwise error:  $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}$

$$\ell(h(\mathbf{x}), f(\mathbf{x}))$$

Defined by the user  
for the ML task!

- Examples:

squared error  $L_2$        $\ell(\hat{y}, y) = (\hat{y} - y)^2$       (regression)

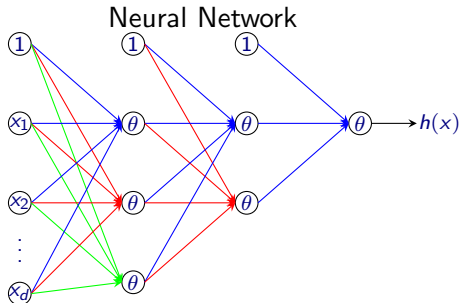
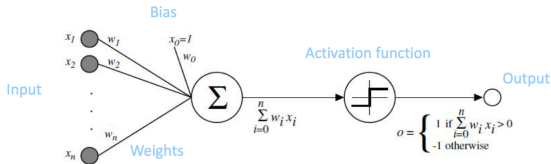
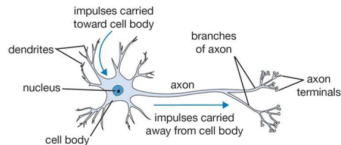
binary error       $\ell(\hat{y}, y) = \mathbb{I}(\hat{y} \neq y)$       (classification)

cross-entropy error       $\ell(\hat{y}, y) = \log(1 + e^{-y_i \mathbf{w}^\top \mathbf{x}_i})$       (see logistic regression)

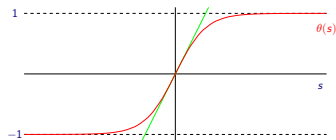
- Training error:  $\hat{R}_n(h) = \frac{1}{n} \sum_{i=1}^n \ell(h(\mathbf{x}_i), y_i)$
- Test error:  $R(h) = \mathbb{E}[\ell(h(\mathbf{x}), y)]$

# Neural Network: Perceptron

- Binary class perceptron
  - biologically inspired model of a single neuron



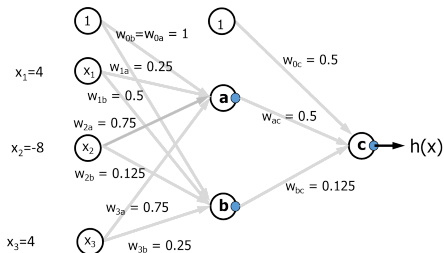
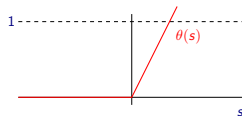
Linear vs. Non-linear activation function  $\theta(s)$



if  $\theta(s) = s$  then  $h(x) = \mathbf{w}_L^\top \mathbf{W}_{L-1} \mathbf{W}_{L-2} \dots \mathbf{W}_1 \mathbf{x} = \mathbf{w}_*^\top \mathbf{x}$

# Neural Network Forward Pass (inference)

- Input  $\mathbf{x} = (4, -8, 4)$
- Non linear activation ReLU  $\theta(s) = s_+ = \max\{0, s\}$ .



- $f = \max(0, \sum_i w_i x_i)$
- $f_a = \max(0, 1 + 0.25 \cdot 4 + 0.75 \cdot (-8) + 0.75 \cdot 4) = 0$
- $f_b = \max(0, 1 + 0.5 \cdot 4 + 0.125 \cdot (-8) + 0.25 \cdot 4) = 3$
- $f_c = h(\mathbf{x}) = \max(0, 0.5 + 0.5 \cdot 0 + 0.125 \cdot 3) = 0.875$
- Ground truth  $y = 2$
- Error  $\hat{R}_n(\mathbf{w}_t) = (y - h(\mathbf{x}))^2 = (2 - 0.875)^2 \approx 1.27$

## Learning Setup with $\mathbf{x} \sim P$ and $P(y|\mathbf{x})$

- Input:  $\mathbf{x} \in \mathcal{X}$
- Output:  $y \in \mathcal{Y}$
- Data:  $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_n, y_n) \sim P$  ( $P$  is the joint distribution of  $(\mathbf{x}, y)$ )

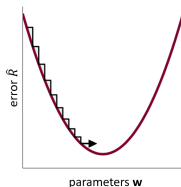
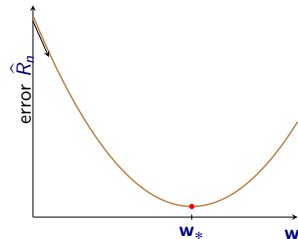
### Learning

- Hypothesis class:  $\mathcal{H} \subset \{h(\mathbf{w}) : \mathcal{X} \rightarrow \mathcal{Y}, w \in \mathbb{R}\}$
- Loss function:  $\ell : \mathcal{Y} \times \mathcal{Y} \rightarrow \mathbb{R}^+$
- Find  $\mathbf{w}^* \in \mathbb{R}^{\|\mathbf{w}\|}$  such that  $g \approx P(y|\mathbf{x})$

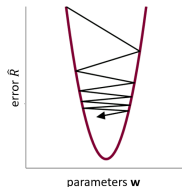
$$\mathbf{w}^* = \underset{\mathbf{w}}{\operatorname{argmin}} \left\{ R(\mathbf{w}) = \mathbb{E} [\ell(h(\mathbf{w}\mathbf{x}), y)] \right\}$$

## Gradient descent

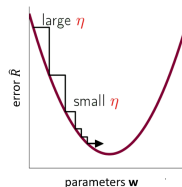
- General method for non-linear optimization:  $\mathbf{w}_{t+1} = \mathbf{w}_t + \eta \mathbf{v}$
- Direction  $\mathbf{v}$ : starting from  $\mathbf{w}_t$ , step along the steepest slope
  - $\mathbf{v} = -\frac{\nabla \hat{R}_n(\mathbf{w}_t)}{\|\nabla \hat{R}_n(\mathbf{w}_t)\|}$  is a unit vector.
- Step size  $\eta$ : how quickly find the minimum
  - $\eta$  is a scalar.



$\eta$  is small



$\eta$  too large



variable  $\eta$

Heuristic: step size should increase with the slope  $\mathbf{w}_{t+1} = \mathbf{w}_t + \Delta \mathbf{w}$

$$\Delta \mathbf{w} = -\eta \nabla \hat{R}_n(\mathbf{w}_t) \text{ (with } \eta \text{ redefined)}$$

# Neural Network Backpropagation (training)

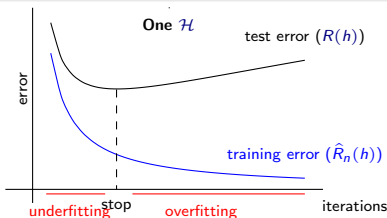
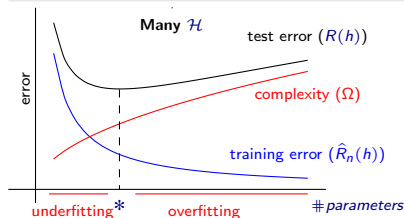
## Backpropagation

- Initialize all weights  $w_{ij}^{(l)}$  at **random**
- for  $t = 1, 2, \dots$  do
  - Pick a data point  $(x_k, y_k)$
  - **Forward:** Compute all  $x_j^{(l)}$
  - **Backward:** Compute all  $\delta_j^{(l)}$
  - **Update:**  $w_{ij}^{(l)} \leftarrow w_{ij}^{(l)} - \eta_t x_i^{(l-1)} \delta_j^{(l)}$ 
    - ★ single point (SGD), minibatch, batch
- Return final weights  $w_{ij}^{(l)}$

# Learning from noisy data

## Overfitting

- Fitting to noise instead of the underlying target function/distribution.
- Occurs when  $\hat{R}_n(h) \downarrow$   $R(h) \uparrow$ , moving away from the target function.



Remember  $n$  matters too!

Noise types: stochastic ( $N \sim \mathcal{N}(0, \sigma^2)$ ),

deterministic (complexity)

deterministic noise	↑
stochastic noise	↑
number of data points	↑

overfitting	↑
overfitting	↑
overfitting	↓

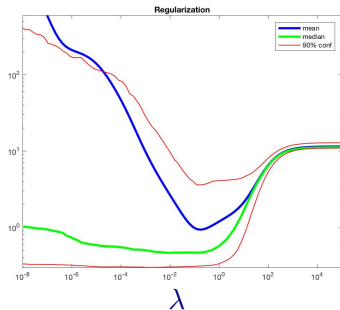
## Regularisation

Regularised Loss (Augmented Error)

$$\mathcal{L}_n(h) = \hat{R}(h) + \lambda \underbrace{\Omega(h)}_{\text{overfit penalty}}$$

•  $\Omega(h)$  – regulariser with parameter  $\lambda$

- $\|\mathbf{w}\|_2^2$ : L2
- $\|\mathbf{w}\|_1$ : L1
- $\|\mathbf{w}\|_1 + \|\mathbf{w}\|_2^2$ : elastic net
- $\mathbf{w}^\top \Gamma^\top \Gamma \mathbf{w}$ : Tikhonov



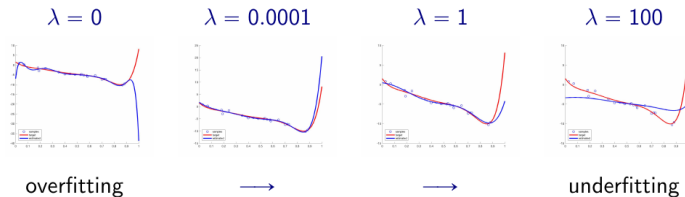
Occam's razor: Simpler is usually better

Regularize towards smoother, simpler functions. Why?

Because noise is not smooth!



# Regularized Loss Minimization



## Regularized Loss Minimization (RLM)

- Hypothesis class  $\mathcal{H} = \cup_i(\mathcal{H}_i, \lambda_i)$ , with  $i \in \mathbb{N}$  e.g.  $\lambda_i \in \{0.0001, 0.001, \dots\}$

- Augmented error:

$$\mathcal{L}_{\bar{k}}(\mathbf{w}, \lambda) = \hat{R}_{\bar{k}}(\mathbf{w}) + \lambda \Omega(\mathbf{w}) \quad \text{e.g. } \Omega(\mathbf{w}) \in \{\|\mathbf{w}\|_2^2, \|\mathbf{w}\|_1, \|\mathbf{w}\|_2^2 + \|\mathbf{w}\|_1, \dots\}$$

- RLM solution:

- for all  $i$ , **train** on  $\mathcal{D}_{\bar{k}}$ :  $g(\mathbf{w}_{\lambda_i}) = \operatorname{argmin}_{\mathbf{w}} \mathcal{L}_{\bar{k}}(\mathbf{w}, \lambda_i)$
- from all  $i$ , **select** on  $\mathcal{D}_k$ :  $g(\mathbf{w}_{\lambda^*}) = \operatorname{argmin}_{\lambda_i} \tilde{R}_k(g(\mathbf{w}_{\lambda_i}))$

## Summary

- Organization of DL course
- Exercises and coursework
- Revision of fundamentals of Machine Learning