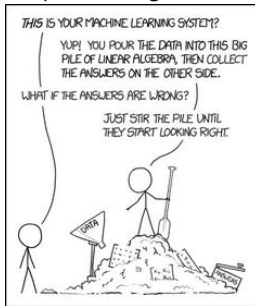


EE3-25: Deep Learning

Krystian Mikolajczyk & Carlo Ciliberto

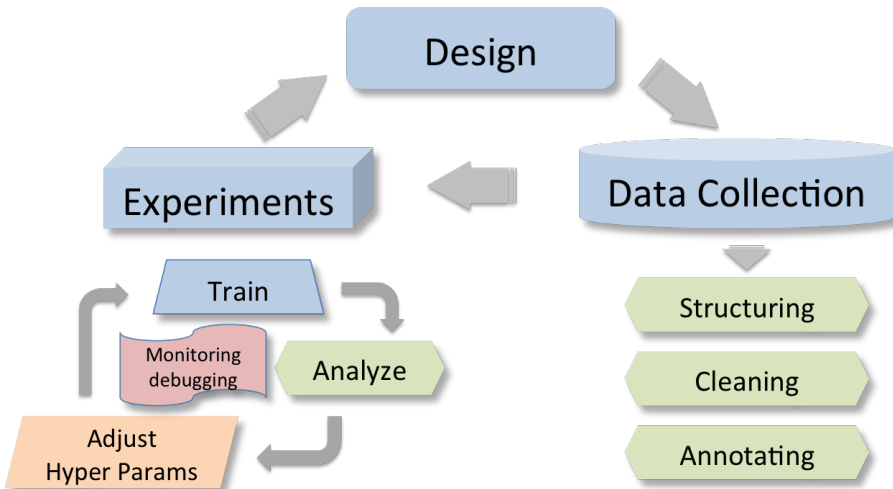
Department of Electrical and Electronic Engineering
Imperial College London



Practical development process

- System design/choice
- Data collection and augmentation
- Hyper parameters search
- Monitoring and debugging

Development process



Design

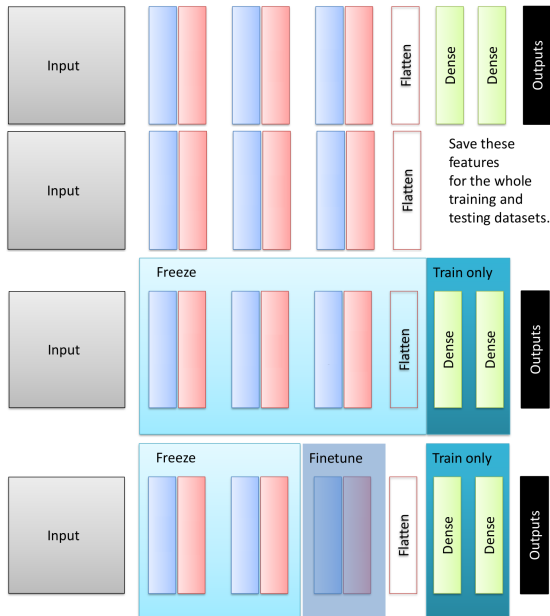
Baseline: end-to-end model

- Was the task studied before? Do literature review!
 - Start from competitions/survey papers
 - Establish a reasonable end-to-end system
- Choose the general category of model based on the structure of your data:
 - MLP for fixed size vectors
 - CNN for images
 - RNN for sequential data
- Nonlinear activations
 - Avoid sigmoid (except for output)
 - ReLU preferred (possibly Leaky ReLU)
 - Use Maxout if most ReLU units die (have zero activation)
- Weights & Biases
 - Random initialization with proper variance
 - For ReLU we prefer a small positive bias to activate ReLU

Design

Finetuning - borrow knowledge

- Pretrain your NN on a large dataset (e.g. same modality, similar task)
 - or start from a pretrained NN
- Option 1: remove / reshape the last few layers and use the features
- Option 2: Fine-tune the parameters on your own dataset
 - Freeze the parameters of first few layers, or make the learning rate small for them
 - Small data - train last FC layers only
 - Medium data - can finetune other layers
 - Use only 1/10th of the original learning rate in finetuning top layer, and 1/100th on intermediate layers



Data Collection

Collect data for the task

- How much data to collect?
 - The more the better
 - Depends on the effect we want to observe
 - Required error bounds and accuracy
- How to label the data?
 - Mechanical Turk, Freelancer, experts, ...
- Avoid bias
 - Selection, Sampling, ...

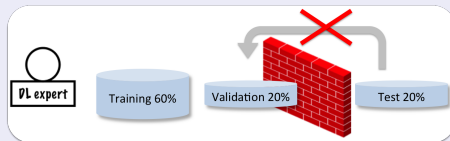
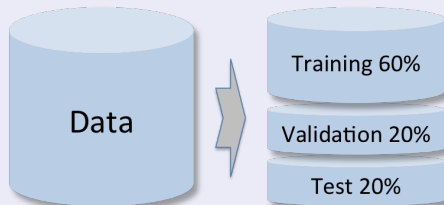
Dataset preparation/curation

- Data structuring & formatting
 - Is the data format suitable?
 - Standardization
- Data cleaning
 - incomplete data, anonymization, missing annotation, correction ...
- Data normalization
 - Value clipping/normalization
 - Whitening

Data Collection

Data split

- Training set
 - typically 60%, to run the learning algorithm on
 - Keep training data balanced
- Validation set
 - typically 20%, to tune hyper parameters, select features
 - make other decisions regarding the learning algorithm
 - also called development set
- Test set
 - typically 20%, to evaluate the performance of the algorithm



Testing

Do not use test data to make any decisions to improve learning!

Validation

Regularisation

$$\mathcal{L}_n(h) = \hat{R}_n(h) + \lambda \underbrace{\Omega(h)}_{\text{overfit penalty}}$$

Validation

$$\underbrace{\mathcal{L}_n(h)}_{\text{direct estimation}} = \hat{R}_n(h) + \lambda \underbrace{\Omega(h)}_{\text{overfit penalty}}$$

- ① Split the training data \mathcal{D} into training \mathcal{D}_{train} and validation \mathcal{D}_{val} sets.
- ② Train g on \mathcal{D}_{train} .
- ③ Estimate its performance on \mathcal{D}_{val} ($v = |\mathcal{D}_{val}|$):

$$\check{R}_v(g) = \frac{1}{v} \sum_{(x_i, y_i) \in \mathcal{D}_{val}} \ell(g(x_i), y_i)$$

Very good estimate of $R(g)$

$$\mathbb{E}_{\mathcal{D}_{val}} [\check{R}_v(g)] \approx R(g) \leq \check{R}_v(g) + \underbrace{\Omega(v, \delta)}_{\sim \sqrt{\log(1/\delta)/v} \leftarrow \text{one model only on } v\text{-points}} \quad \text{w. p. } 1 - \delta$$

- \mathcal{D}_{val} is unbiased, small Hoeffding bound, only one g is considered.
- Select $\lambda^* = \operatorname{argmin}_{\lambda} \check{R}_v(g)$, then train on the whole \mathcal{D} with λ^* .

Validation: More generally

Validation

Given hypothesis classes $(\mathcal{H}_1, \lambda_1), \dots, (\mathcal{H}_i, \lambda_i), \dots, (\mathcal{H}_M, \lambda_M)$,

- ① Split training data \mathcal{D} into \mathcal{D}_{train} and \mathcal{D}_{val} sets.
- ② Train g_i on \mathcal{D}_{train} , $|\mathcal{D}_{train}| = n - v$
- ③ Select i_* such that $g_{i_*} = \operatorname{argmin}_i \check{R}_v(g_i)$ on $|\mathcal{D}_{val}| = v$
 - i - not only regularisation, can be other hyperparameter.
- ④ Select $\mathcal{H}_{i_*}, \lambda_{i_*}$ and train new g on $\mathcal{D}_{val} \cup \mathcal{D}_{train}$ to get the final g_* .

- Cost? $(n - v)$ and learning curves.
- if $v \uparrow$ then $\check{R}_v(g) \sim R(g) \uparrow$, but $|\mathcal{D}_{train}| = n - v \downarrow$ then $R(g) \uparrow$
- How big $|\mathcal{D}_{val}| = \frac{|\mathcal{D}_{train}|}{5}$
- Why “validation” and not “test” set? Unlucky split of dataset \mathcal{D} ?

Validation

Cross validation

- ① Split data \mathcal{D} into K disjunct parts: $\mathcal{D} = \cup_{k=1}^K \mathcal{D}_k$
 - ② For each k , create training and validation set:
 - Training $\mathcal{D}_{\bar{k}} = \mathcal{D}_1 \cup \mathcal{D}_2 \cup \dots \cancel{\mathcal{D}_k} \cup \dots \cup \mathcal{D}_K$, $|\mathcal{D}_{\bar{k}}| = \bar{n} = \frac{K-1}{K}n$
 - Validation \mathcal{D}_k , $|\mathcal{D}_k| = \bar{k} = \frac{n}{K}$
 - ③ Train $g_{\bar{k}}$ on $\mathcal{D}_{\bar{k}}$ with ERM
- Validation error of $g_{\bar{k}}$ on \mathcal{D}_k is $\check{R}_k(g_{\bar{k}}) = \frac{1}{|\mathcal{D}_k|} \sum_{(\mathbf{x}_i, y_i) \in \mathcal{D}_k} \ell(g_{\bar{k}}(\mathbf{x}_i), y_i)$
 - Cross-validation error – K -fold cross validation

$$\check{R}_{CV} = \frac{1}{K} \sum_{k=1}^K \check{R}_k(g_{\bar{k}})$$

Cross validation: Is this good?

Expected performance of the learnt hypothesis: $\mathbb{E}_{\mathcal{D}} [R(g^{(\mathcal{D})})]$

$$\begin{aligned}\mathbb{E}_{\mathcal{D}_k} [\check{R}_k(g_{\bar{k}})] &= \mathbb{E}_{\mathcal{D}_k} \left[\frac{1}{|\mathcal{D}_k|} \sum_{(\mathbf{x}_i, y_i) \in \mathcal{D}_k} \ell(g_{\bar{k}}(\mathbf{x}_i), y_i) \right] \\ &= \mathbb{E}_{\mathbf{x}, y} [\ell(g_{\bar{k}}(\mathbf{x}), y)] \\ &\approx R(g_{\bar{k}}) = R(g^{(\mathcal{D}_{\bar{k}})})\end{aligned}$$

If all chunks are of the same size, then $g_{\bar{k}}$ is trained on $n = \frac{K-1}{K}n$ points:

$$\mathbb{E}_{\mathcal{D}} [\check{R}_k(g_{\bar{k}})] \approx \mathbb{E}_{\mathcal{D}_k} [R(g_{\bar{k}})] \approx \mathbb{E}_{\mathcal{D}} [\check{R}_{CV}]$$

Cross validation

\mathcal{D}_1	\mathcal{D}_2	\mathcal{D}_3	\cdots	\mathcal{D}_K	$\rightarrow \check{R}_1(g_{\bar{1}})$
\mathcal{D}_1	\mathcal{D}_2	\mathcal{D}_3	\cdots	\mathcal{D}_K	$\rightarrow \check{R}_2(g_{\bar{2}})$
\mathcal{D}_1	\mathcal{D}_2	\mathcal{D}_3	\cdots	\mathcal{D}_K	$\rightarrow \check{R}_3(g_{\bar{3}})$
\vdots	\vdots	\vdots	\ddots	\vdots	
\mathcal{D}_1	\mathcal{D}_2	\mathcal{D}_3	\cdots	\mathcal{D}_K	$\rightarrow \check{R}_K(g_{\bar{K}})$

Cross-validation error

$$\check{R}_{CV} = \frac{1}{K} \sum_{k=1}^K \check{R}_k(g_{\bar{k}})$$

$$\mathbb{E}_{\mathcal{D}} [\check{R}_{CV}] \approx \mathbb{E}_{\mathcal{D}_k} [R(g_{\bar{k}})]$$

with $\bar{n} = (K - 1/K)n$:

Typical K choices:

- $K = 10$: 10-fold cross validation
- $K = n$: leave-one-out cross validation

Cross validation

K-fold Cross Validation

Given: hypothesis classes $(\mathcal{H}_1, \theta_1), \dots, (\mathcal{H}_i, \theta_i), \dots, (\mathcal{H}_M, \theta_M)$,

with θ – any hyperparameter

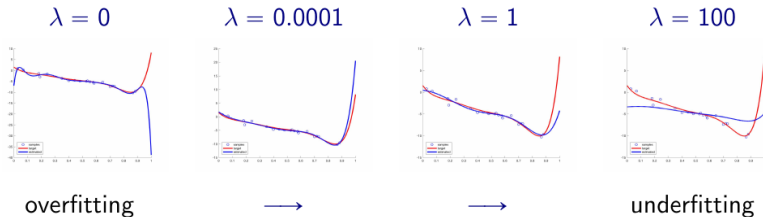
K training sets $\mathcal{D}_{\bar{k}} = \mathcal{D}_1 \cup \mathcal{D}_2 \cup \dots \setminus \mathcal{D}_k \cup \dots \cup \mathcal{D}_K$, $|\mathcal{D}_{\bar{k}}| = \bar{n} = \frac{K-1}{K}n$

K validation sets \mathcal{D}_k , $|\mathcal{D}_k| = \bar{k} = \frac{n}{K}$

- 1 Train $g_{i,\bar{k}}$ with ERM for every $(\mathcal{H}_i, \theta_i)$ and every $\mathcal{D}_{\bar{k}}$
- 2 Compute $\check{R}_{CV}(g_i)$ for every g_i on all \mathcal{D}_k
- 3 Select $g_{i_*} = \operatorname{argmin}_{g_i} \check{R}_{CV}(g_i)$
- 4 Use $(\mathcal{H}_{i_*}, \theta_{i_*})$ and whole data set \mathcal{D} to train final g_*

Regularized Loss Minimization

Regularization



Validation

\mathcal{D}_1	\mathcal{D}_2	\mathcal{D}_3	\dots	\mathcal{D}_K	$\rightarrow \check{R}_1(g_{\bar{1}})$
\mathcal{D}_1	\mathcal{D}_2	\mathcal{D}_3	\dots	\mathcal{D}_K	$\rightarrow \check{R}_2(g_{\bar{2}})$
\mathcal{D}_1	\mathcal{D}_2	\mathcal{D}_3	\dots	\mathcal{D}_K	$\rightarrow \check{R}_3(g_{\bar{3}})$
\vdots	\vdots	\vdots	\ddots	\vdots	
\mathcal{D}_1	\mathcal{D}_2	\mathcal{D}_3	\dots	\mathcal{D}_K	$\rightarrow \check{R}_K(g_{\bar{K}})$

$$\check{R}_{CV} = \frac{1}{K} \sum_{k=1}^K \check{R}_k(g_{\bar{k}})$$

Regularized Loss Minimization (RLM)

- Hypothesis class $\mathcal{H} = \cup_i(\mathcal{H}_i, \lambda_i)$, with $i \in \mathbb{N}$ e.g. $\lambda_i \in \{0.0001, 0.001, \dots\}$

- Augmented error:

$$\mathcal{L}_{\bar{k}}(\mathbf{w}, \lambda) = \hat{R}_{\bar{k}}(\mathbf{w}) + \lambda \Omega(\mathbf{w})$$

e.g. $\Omega(\mathbf{w}) \in \{\|\mathbf{w}\|_1, \|\mathbf{w}\|_2^2, \|\mathbf{w}\|_Q^2, \dots\}$

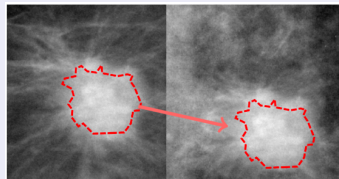
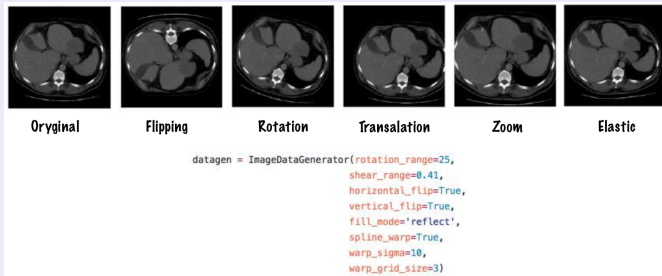
- RLM solution:

- for all i , train on $\mathcal{D}_{\bar{k}}$: $g(\mathbf{w}_{\lambda_i}) = \operatorname{argmin}_{\mathbf{w}} \mathcal{L}_{\bar{k}}(\mathbf{w}, \lambda_i)$
- from all i , select on \mathcal{D}_k : $g(\mathbf{w}_{\lambda^*}) = \operatorname{argmin}_{\lambda_i} \check{R}_k(g(\mathbf{w}_{\lambda_i}))$

Data Augmentation

Image Data Augmentation

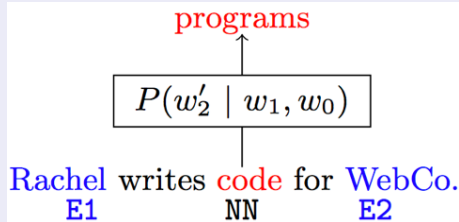
- Adding noise
- Generating modified samples
- Medical data
 - Segment tumor mass
 - Move
 - Resample background tissue
 - Blend



Data Augmentation

Text Data Augmentation

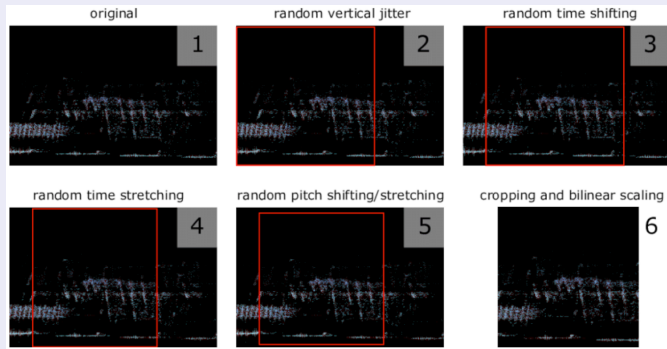
- Adding noise
- Inserting synonyms
- A conditional word-swap with externally trained language model and specifically targeting nouns (NN) between entity mentions (E1,E2)
- Rare words in new, synthetically created contexts



Data Augmentation

Audio Data Augmentation

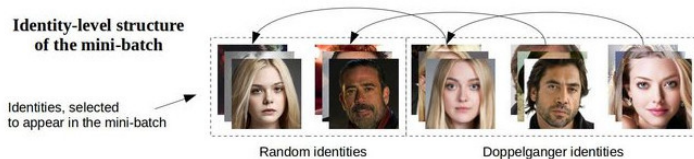
- Adding noise
- Vertical jitter
- Time shifting
- Time stretching: change the speed of the audio signal
- Pitch shift
- Cropping and bilinear scaling



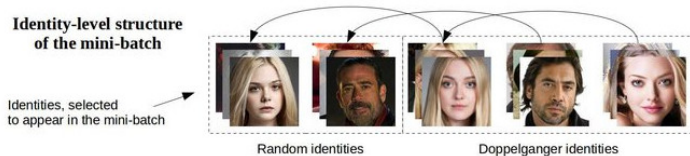
Hard Positive/Negative Mining

Adversarial training

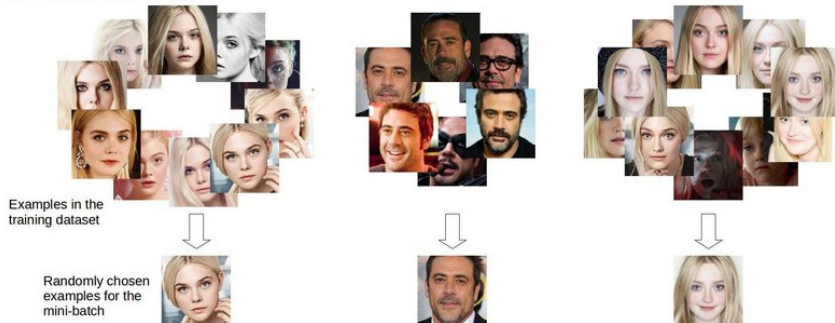
- Train the system from randomly formed mini-batches with balanced positive/negative examples
- Identify hard examples (close to decision boundary) during validation
- Form a new mini-batch by including the hard examples from the the previous iteration.



Hard Positive/Negative Mining

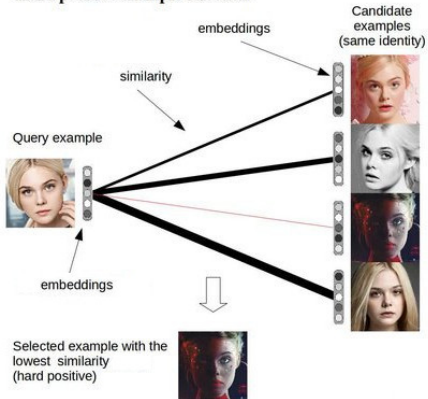


Random example selection

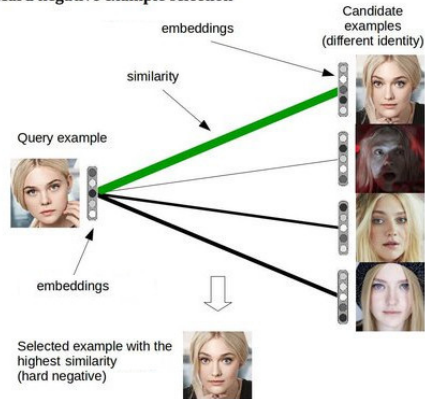


Hard Positive/Negative Mining

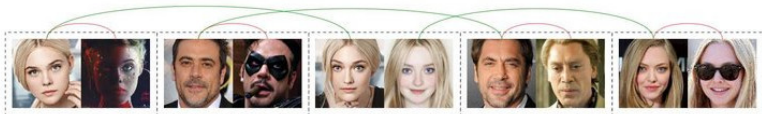
Hard positive example selection



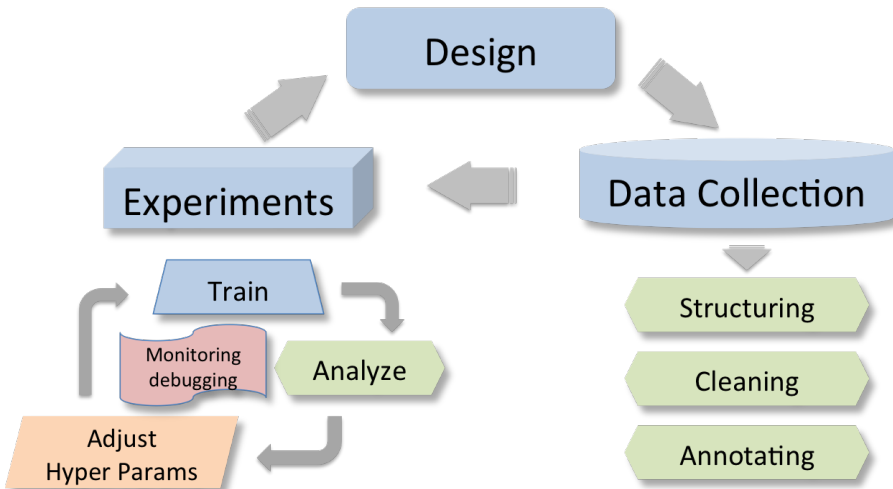
Hard negative example selection



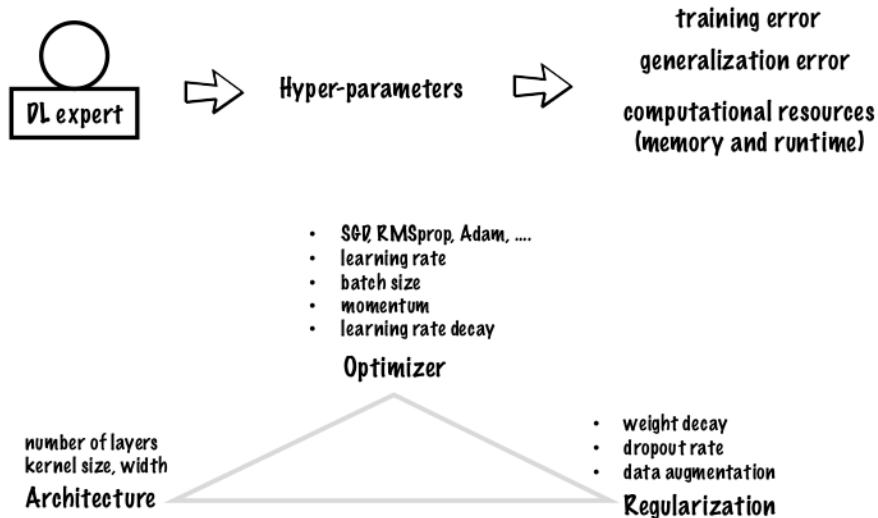
Example-level structure of the mini-batch



Development process



Parameter setting



Parameter setting

Hyper-parameters

- Optimization of hyper-parameters affects the quality of local minima that our model can reach (effective model capacity)
 - memory problems? reduce batch size
- The bigger the architecture the higher the model's capacity
 - caveat: memory and computation time are limited
- To decrease generalization gap increase regularization
 - increase weight decay and dropout rate to reduce model capacity
 - data augmentation does not affect model capacity

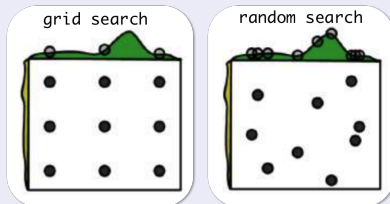
Good rule of thumb

Select the default parameters!

Parameter setting

Hyper parameter search: Good intuition and experience

- Learning Rate / Momentum
 - Decrease learning rate while training
 - Typical momentum to 0.8 - 0.9
- Batch Size
 - Shuffle data before batching
 - For large dataset: set to whatever fits your memory
 - For smaller dataset: find a tradeoff between instance randomness and gradient smoothness
- More efficient to optimize hyper-parameter with randomly chosen trials rather than on grid
- Model based hyper-parameter selection
- Use coarse to fine search for hyper-parameters
- Search on log scale (eg. learning rates)



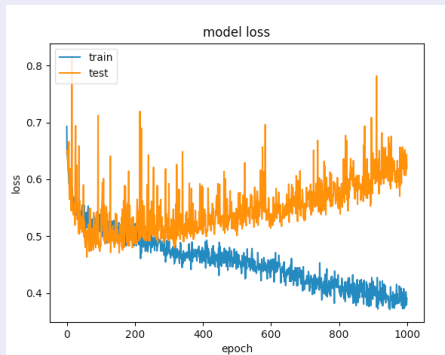
Monitoring and debugging

- Set a small dataset
- Reasoning using train and validation error/plots
- Monitoring histograms of activations and gradient updates
- Analysing of model predictions and errors

Monitoring and debugging

Set a small dataset

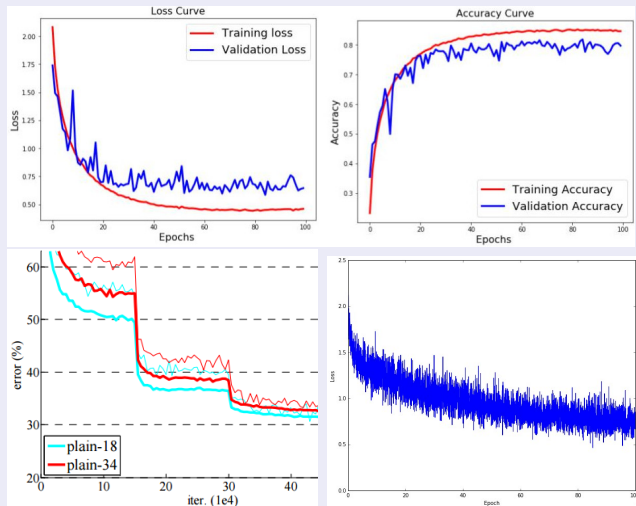
- Observe training error/plot
 - The model should overfit
- Change the optimisation parameters if the model does not overfit
 - If unsuccessful, inspect the code for potential bugs
- Reduce the size of the model (architecture) if model compilation time is long



Monitoring and debugging

- Optimize your hyper-parameter in validation and evaluate on test
- Keep track of training and validation loss during training
- Do early stopping if training and validation loss diverge
- Use patience - wait this long from last change
- Relative improvement threshold (significance)
- Loss does not tell you all. Try precision, class-wise precision, and other metrics

Train and validation error/plots



Monitoring and debugging

Visualize model predictions

- Check input data and annotations
- Check if predictions make sense

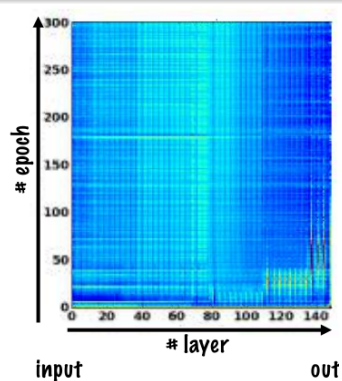
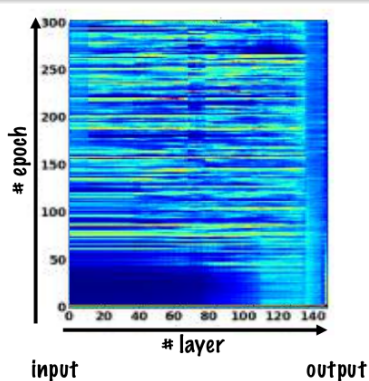
Understanding the underlying causes of the errors

- Get a sample of e.g. 100 validation set examples for which the system failed
- Examine these examples manually/visually
- Identify the most common errors
 - A simple excel table might be enough

Monitoring and debugging

Monitor histograms of activations and gradient updates per layer

- Erratic/unstable
- Uniform



Monitoring and debugging

Generalisation problem

- Biased data distributions (not i.i.d.)

Training & validation data



Test data



- Mix datasets and test
- Split validation into training validation and test validation

Training & validation data

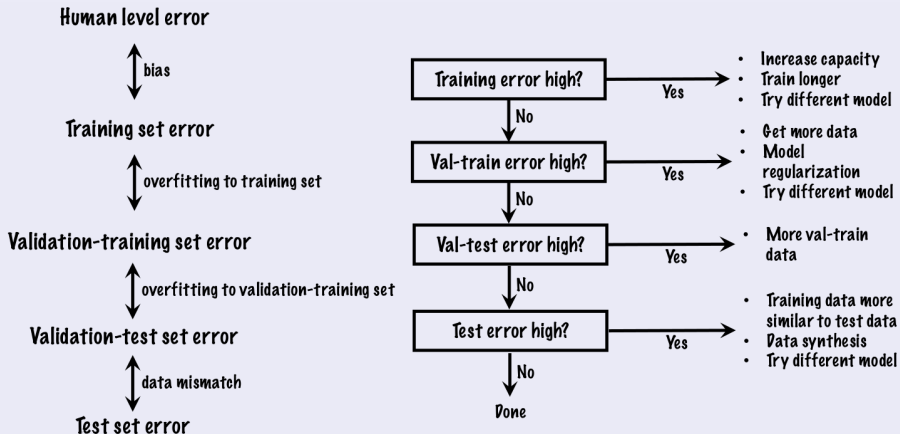


Test & validation data



Monitoring and debugging

Error analysis



Summary

- System design/choice
- Data collection and augmentation
- Hyper parameters search
- Monitoring and debugging